Governing with AI, Governing AI: Public Innovation and Democratic Resilience in the United States and Taiwan

Karl Ho
University of Texas at Dallas
karlho@gmail.com

Peter C.H. Pan
National Chung Hsing University
chpan@nchu.edu.tw

Abstract

Public agencies in the United States and Taiwan are adopting artificial intelligence under new governance regimes, but adoption alone does not establish administrative capacity or democratic legitimacy. This chapter asks how the arrangements surrounding public-sector AI shape whether agencies remain answerable for it. We distinguish *governing with AI*, the use of AI in services, decision support and communication, from *governing AI*, the authorization, procurement, risk assessment, disclosure and redress mechanisms through which agencies oversee those systems. We develop the concept of accountable adaptive capacity: the ability of public institutions to experiment with AI, evaluate its effects, correct failures and sustain accessible services while keeping responsibility identifiable and open to challenge. The United States governs federal AI through executive memoranda that changed direction in 2025; Taiwan's AI Basic Act, in force since January 2026, distributes implementation across ministries under a new national strategy committee. The design combines a coded analysis of the US federal AI use-case inventories, a parallel coding of Taiwanese agency rules and procurement records, and four matched cases. Preliminary evidence from the 2025 US inventory shows a wide gap between formal requirements and reported implementation: of 227 deployed high-impact uses, fewer than one in five report a completed impact assessment, and 7 percent report an established appeal process. Throughout, we distinguish formal requirements, agency-reported compliance and independently supported implementation.

Adoption Without Answerability

Public agencies in both the United States and Taiwan now use artificial intelligence at scale. The US federal inventory of AI use cases grew from 2,133 entries in 2024 to 3,611 in 2025 (Office of Management and Budget, 2025a, 2026). Taiwan's cabinet authorized civil servants to use generative AI in 2023 and, under the AI Basic Act that took effect in January 2026, ordered every agency to complete an AI-use risk assessment by July 2026 (Executive Yuan, 2023; Legislative Yuan, 2026; Executive Yuan, 2026).[1]

Both governments pair adoption with commitments to trust. Neither has yet shown that its agencies can answer for the systems they deploy: that someone identifiable owns each system, that its effects are evaluated, that errors are corrected, and that the people it affects can challenge its outputs. This gap between adoption and answerability is the subject of this chapter.

We make three contributions. First, we separate two tasks that policy language routinely merges. *Governing with AI* concerns the use of AI in services, decision support, enforcement and communication. *Governing AI* concerns the authorization, procurement, risk assessment, disclosure and redress arrangements through which public institutions oversee those uses. Second, we propose *accountable adaptive capacity* as the quantity that links the two, and specify it in observable terms. Third, we compare a regime of executive memoranda with a regime of framework legislation, using coded administrative records rather than policy texts alone.

[Draft note: the introduction to the volume (Chapter 1) sets out the general argument; this chapter should not restate it. Keep this section to about 700 words.]

Accountable Adaptive Capacity

Two meanings of governing

The public administration literature has long recognized that information technology changes where discretion sits. Bovens and Zouridis (2002) described the move from street-level to system-level bureaucracy, in which decisions once made by front-line officials are embedded in system design. Work on artificial discretion extends the point to machine learning (Young et al., 2019; Bullock, 2019). A second literature asks how algorithmic systems can be held to account (Busuioc, 2021; Wieringa, 2020; Raji et al., 2020), building on the general account of accountability as a relationship in which an actor must explain and justify

[1]This paper is in the first draft version, prepared for the International Conference on Public Governance and Practice Forum, National Chung Hsing University, November 2026. The authors thank the conference organizers and participants for their feedback. It serves as the comparative chapter of the edited volume *Governing with AI, Governing AI: Public Innovation and Democratic Resilience*.

conduct to a forum that can pose questions and impose consequences (Bovens, 2007).

The two literatures map onto the two meanings of governing. Studies of discretion are mostly about governing *with* AI: what happens to decisions when an algorithm shapes them (Alon-Barkat and Busuioc, 2023; Selten et al., 2023). Studies of accountability are mostly about governing AI: what arrangements make a system answerable (Grimmelikhuijsen and Meijer, 2022; Yeung, 2018). Public value scholarship sits between them (Moore, 1995; Chen et al., 2023; Schiff et al., 2022). We argue that the democratic consequences of public-sector AI depend on how the two are coupled.

A third body of work studies adoption itself. Reviews document rapid diffusion and persistent governance gaps (Wirtz et al., 2019; Zuiderwijk et al., 2021; Madan and Ashok, 2023), and comparative case studies show that organizational capacity shapes how agencies adopt AI (Neumann et al., 2024; Kuziemski and Misuraca, 2020). For the United States, two empirical assessments anchor our analysis: Engstrom et al. (2020) mapped AI use across federal agencies, and Lawrence et al. (2023) found that fewer than 40 percent of 45 legal requirements could be verified as implemented and that nearly half of agencies had not published AI use-case inventories. The Government Accountability Office likewise found that agencies had begun but not completed key requirements (U.S. Government Accountability Office, 2023). We extend this line of work in two directions: to the post-2025 regime, whose inventory reports safeguard status use by use, and to a second democracy whose governance runs through framework legislation rather than executive memoranda. Taiwan also brings a distinct tradition of participatory digital governance (Hsiao et al., 2018), which makes the question of public challenge especially sharp.

Four components

We define accountable adaptive capacity as the ability of public institutions to experiment with AI, evaluate its effects, correct failures and sustain accessible services while keeping responsibility for decisions identifiable and open to public challenge. It has four components, each with observable indicators.

Table 1: Components of accountable adaptive capacity and their indicators.

Component	What it requires	Indicators in administrative records
Experimentation	Controlled pilots with defined scope and stopping rules	Pilot stage recorded; pre-deployment testing
Evaluation	Evidence on effects, not only on adoption	Impact assessment; ongoing monitoring; independent review

Component	What it requires	Indicators in administrative records
Correction	Routes to fix, suspend or withdraw a system	Fail-safe or human override; retirement records; termination of non-compliant uses
Answerability	An identifiable owner and an accessible channel for challenge	Named owner; notice to affected people; appeal or remedy process

The concept draws on experimentalist governance, in which provisional goals are revised through monitored local experimentation (Sabel and Zeitlin, 2012), and on work treating public-sector resilience as the capacity to absorb disruption while sustaining core functions (Boin and van Eeten, 2013; Duit, 2016; Ansell et al., 2023). We add the democratic condition: resilience that preserves services but loses identifiable responsibility, equal treatment or contestability is organizational resilience, not democratic resilience (Merkel and Lührmann, 2021; Boese et al., 2021).

Propositions

Table 2: Propositions.

ID	Proposition	Observable implication
P1 Coupling	Safeguards are more complete where use has an identifiable owner and a documented correction route	Deployed high-impact uses with named owners report more established minimum practices
P2 Procurement	Vendor-supplied systems are less inspectable than in-house systems	Purchased systems report fewer established evaluation and audit practices, net of agency
P3 Classification discretion	Agencies use definitional discretion to narrow the scope of heightened oversight	Uses rated rights- or safety-impacting in 2024 are classified out of high-impact status in 2025
P4 Distributed implementation	Under Taiwan's sectoral model, ministries with established digital units adopt fuller answerability provisions	Variation in agency rules tracks prior digital capacity

ID	Proposition	Observable implication
P5 Transfer	Routines transfer between systems more readily than statutory instruments	Staff review and fallback routines appear in both systems; legal instruments do not

Two Regimes

[Draft section, about 1,400 words. Structure the comparison on the five dimensions used throughout the volume. One paragraph per dimension, each ending with what the difference implies for accountable adaptive capacity.]

Authority and responsibility. In the United States, federal AI governance rests on executive memoranda. M-24-10 established minimum practices for rights- and safety-impacting AI in 2024 (Office of Management and Budget, 2024); M-25-21 replaced it in April 2025 with a single high-impact category, Chief AI Officers in every agency and a one-year compliance window (Office of Management and Budget, 2025b). The NIST AI Risk Management Framework remains the main non-binding technical reference (National Institute of Standards and Technology, 2023). In Taiwan, the AI Basic Act names the National Science and Technology Council as central competent authority, tasks the Ministry of Digital Affairs with a risk-classification framework, and leaves sector rules to line ministries (Legislative Yuan, 2026). A cabinet-level strategy committee chaired by the premier coordinates implementation (Executive Yuan, 2026).

Procurement and data. M-25-22 governs federal AI acquisition, including vendor lock-in and data rights (Office of Management and Budget, 2025c); M-26-04 adds contract terms on model neutrality (Office of Management and Budget, 2025d). *[Taiwan: procurement law and any AI-specific contract guidance; check MODA and Public Construction Commission sources.]*

Workforce capability. *[US: CAIO role and AI workforce provisions in M-25-21. Taiwan: civil-service AI competency guidelines announced for June 2026 by the Directorate-General of Personnel Administration and MODA.]*

Transparency and public challenge. The US relies on annual public inventories and public reporting of high-impact determinations and waivers (Office of Management and Budget, 2025b). Taiwan's 2023 guidelines require disclosure when agencies use generative AI in service delivery, and prohibit unverified AI output as the sole basis for administrative decisions (Executive Yuan, 2023). *[Add: whether Taiwan agencies' July 2026 risk assessments are public.]*

Learning and service continuity. *[Fail-safe and fallback provisions in each regime; Taiwan's internal control rules due within one year of May 2026.]*

Data and Design

Evidence and its limits

The two jurisdictions do not produce symmetric evidence. The United States publishes a machine-readable inventory of individual AI uses with fields that map directly onto the minimum practices in M-25-21. Taiwan publishes no equivalent, so the Taiwanese side must be built from agency rules, procurement records and information requests. We accept this asymmetry and assign each component of the design a distinct task. The US inventory tests P1–P3 at scale. Procurement records, which both governments publish, provide the symmetric comparison. Agency rules and matched cases carry P4 and P5.

Throughout, we distinguish three levels of evidence: what a regime requires, what agencies report about their own compliance, and what independent sources support. Inventory fields are the second kind. They measure disclosed governance, not practice.

Datasets

Table 3: Datasets for the comparison.

Dataset	Coverage	Unit	Use	Status
Federal AI Use Case Inventory, 2024 and 2025 (Office of Management and Budget, 2025a, 2026)	US federal	Use case (2,133; 3,611)	P1–P3	Downloaded
Agency AI compliance plans and strategies under M-25-21	US federal	Agency	Capacity covariates	To collect
Federal contract awards (USAspending.gov / FPDS)	US federal	Contract	Symmetric procurement comparison	To collect

Dataset	Coverage	Unit	Use	Status
Government e-Procurement System award notices (; data.gov.tw)	Taiwan central and local	Tender	Symmetric procurement comparison	To collect
Ministry internal AI rules and risk-assessment summaries	Taiwan central	Agency (about 30)	P4	Information requests
Executive Yuan guidelines, AI Basic Act, implementing directives	Taiwan	Document	Requirements coding	Collected
Government AI Readiness Index (Oxford Insights, 2024)	Both, cross-national	Country-year	Context only	Public
Public opinion on government AI (Huang, 2025); Pew Research Center surveys	Both	Respondent	Legitimacy context; item comparability	Public; items differ
Matched-case documents and interviews	Both	Case (4)	P5	IRB pending

Two cross-national resources are of limited use. The UN E-Government Survey does not cover Taiwan, and most cross-national AI indices score national ecosystems rather than public-sector governance. The Government AI Readiness Index covers both countries; Taiwan ranked fourth in East Asia in 2024 (Oxford Insights, 2024). We use it only to situate the cases.

Data collection plan

1. **US inventory (complete for 2024–2025).** Link 2024 and 2025 records. Exact name matching within agency recovers 1,075 of the 3,611 records in 2025, about 30 percent. Raise the match rate with fuzzy string matching on name and bureau, then adjudicate uncertain pairs by hand.
2. **Procurement, both countries (December 2026 – February 2027).** Retrieve award records for 2023–2026 by keyword (, AI, , , ; artificial intelligence, machine learning, generative, large language model) and classify relevance with an LLM, validated against a hand-coded sample. Code procurement route, vendor, value, agency and any evaluation, audit or data-rights terms in the tender documents.
3. **Taiwan agency rules (December 2026 – March 2027).** Collect posted internal AI rules from about 30 central agencies; file requests under the Freedom of Government Information Law for the July 2026 risk assessments where they are not posted.
4. **Matched cases (January – March 2027).** Two functions in each country, matched on task, affected population and decision stakes: a public-facing generative AI assistant, and AI triage of claims or inspections with human final decision. Candidate US cases come from the inventory; Taiwanese candidates from agency disclosures and procurement records.
5. **Interviews (subject to UTD IRB approval).** Four to six per case with system owners, oversight staff and vendors.
6. **Optional: parallel survey experiment.** Field identical vignettes in both countries varying human review, appeal, disclosure and vendor provenance, with perceived legitimacy as the outcome. This addresses the non-comparability of existing opinion items.

Method and model route

The design nests four components. Each answers a different question, so no single estimate carries the comparison.

A. Common codebook. Every use case, agency rule and case is coded on the four components in Table 1 and the volume's five dimensions, each scored 0 (absent), 1 (stated) or 2 (documented as implemented). Free-text fields and Chinese-language documents receive a first-pass coding by a large language model (Gilardi et al., 2023), with two human coders on a 20 percent stratified sample. Items below a Krippendorff's alpha of 0.80 are recoded (Krippendorff, 2018).

B. US inventory model (P1–P3). For each deployed high-impact use case i in agency j, let y_{ij} be the number of M-25-21 minimum practices reported as established out of the n_{ij} that apply. We fit a Bayesian multilevel binomial model (Gelman and Hill, 2007; Bürkner, 2017):

$$y_{ij} \sim \text{Binomial}(n_{ij}, \pi_{ij}),$$

$$\text{logit}(\pi_{ij}) = \alpha + \beta_1 \text{Vendor}_{ij} + \beta_2 \text{Mixed}_{ij} + \beta_3 \text{GenAI}_{ij} + \beta_4 \text{PII}_{ij} + \gamma \, \text{Capacity}_j + u_j,$$

with $u_j \sim N(0, \sigma_u^2)$. Agency intercepts are essential: the descriptive association between vendor procurement and incomplete safeguards (below) is largely an agency composition effect. An ordinal specification on the 0–2 codebook score is a robustness check (Bürkner and Vuorre, 2019). Uses that report none of the fields are treated as missing outcomes, not as zero safeguards, with a sensitivity analysis that bounds the estimates under both assumptions. For P3, a multilevel logit models 2025 high-impact status among uses rated rights- or safety-impacting in 2024.

C. Procurement comparison (symmetric). Monthly counts of AI-related awards in each country are modeled as interrupted time series around regime events (M-25-21 in April 2025; the AI Basic Act in January 2026), with contract-level models of evaluation and audit terms by procurement route.

D. Taiwan corpus and matched cases (P4, P5). With about 30 agencies, the Taiwan corpus supports descriptive comparison and simple ordinal models relating rule completeness to prior digital capacity. The four cases are analyzed by structured, focused comparison with process tracing of one authorization-to-correction chain per case (George and Bennett, 2005; Beach and Pedersen, 2019). Results are reported as a case-by-dimension matrix with the evidence level marked in each cell.

Integration. Component B places each US case in the inventory distribution, so we know whether a case is typical. Transfer claims are drawn only from C and D, where both countries are observed on equal terms (Dolowitz and Marsh, 2000).

Preliminary Evidence: The US Inventory

Concentration and composition

Table 4: US Federal AI Use Case Inventory, 2025: summary.

Measure	Value
Deployed high-impact (2025)	227
Use cases (2024)	2,133
Use cases (2025)	3,611
Reporting agencies (2025)	41
High-impact (2025)	445
Presumed high-impact, determined not (2025)	110
Deployed or pilot (2025)	1,480

Table 4: US Federal AI Use Case Inventory, 2025: summary.

Measure	Value
Generative or agentic AI (2025)	975
Vendor purchased (2025)	594
In-house (2025)	581
Contracting and in-house (2025)	429

High-impact AI is concentrated in three departments. The Department of Veterans Affairs reports 215 of the 445 high-impact uses, the Department of Justice 114 and Homeland Security 55 (Figure 1). Any government-wide average is therefore largely a statement about these three agencies.

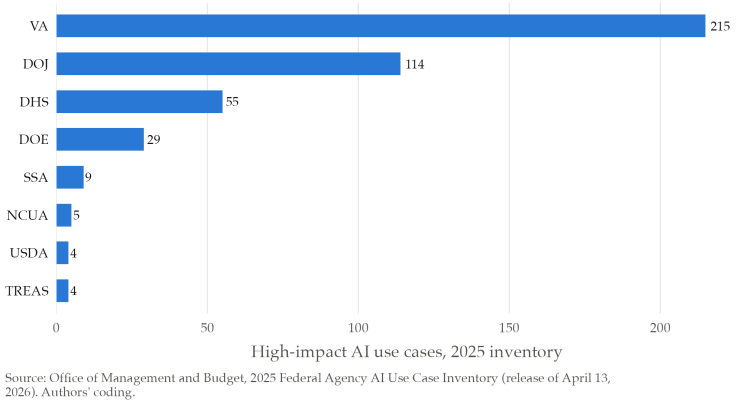
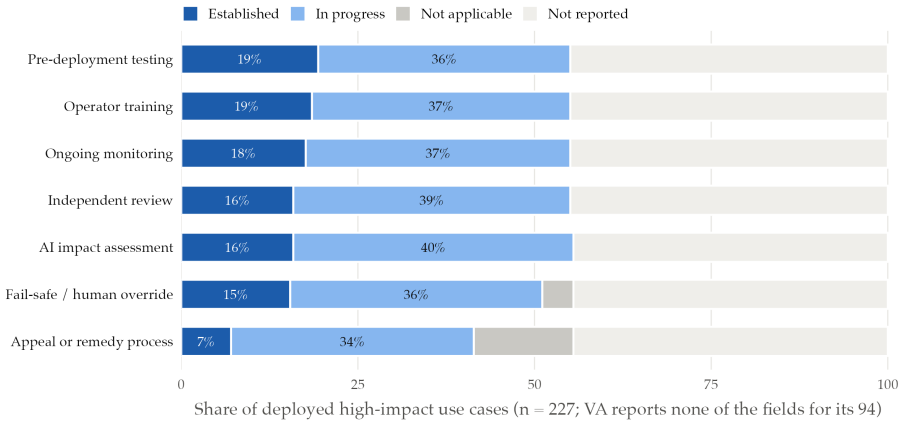


Figure 1: High-impact AI use cases by agency, 2025 inventory.

The implementation gap

OMB requires the minimum-practice fields only for deployed high-impact uses, so the relevant population is the 227 such uses. Figure 2 reports their status. For every practice, the modal reported status is "in progress"; fewer than one in five report an established practice, and appeal or remedy processes are established for 7 percent. The Department of Veterans Affairs reports none of these fields for its 94 deployed high-impact uses.

[Interpretation: the inventory was compiled during the 365-day compliance window, which closed in April 2026, so "in progress" is the expected status for many uses. The 2026 inventory will show whether these commitments were completed. The appeal gap is the substantive finding: it is the practice most directly tied to contestability, and it lags the others.]



Source: Office of Management and Budget, 2025 Federal Agency AI Use Case Inventory (release of April 13, 2026). Authors' coding.

Figure 2: Status of M-25-21 minimum practices among deployed high-impact AI uses, 2025 inventory. "Not reported" is dominated by the Department of Veterans Affairs.

Table 5: Mean share of applicable minimum practices established, deployed high-impact uses excluding the Department of Veterans Affairs, by procurement route.

Route	n	Mean completeness	Share with any practice established
Contracting and In House	30	0.64	0.77
In-house Development	16	0.62	0.62
Vendor Purchased	84	0.14	0.18

Vendor-purchased systems report far fewer established practices than in-house or mixed systems (Table 5). This is suggestive for P2 but not yet evidence for it: 65 of the 84 vendor-purchased uses belong to the Department of Justice, which reports every practice as in progress. The multilevel model in Component B separates the procurement effect from the agency effect.

Reclassification under the new definition

Table 6: 2025 determination of uses matched to the 2024 inventory, by 2024 risk category (exact-name matches).

2024 category	High-impact	Not High-impact	Not reported	Presumed High-Impact, but Not High-impact
Neither	8	731	77	34
Other / too new	3	5	0	1
Rights and/or safety-impacting	180	22	0	14

Of the 216 matched uses rated rights- or safety-impacting in 2024, 180 (83 percent) are high-impact in 2025; 36 (17 percent) left heightened oversight, 14 of them through the new "presumed high-impact but determined not" category (Figure 3). The matched set covers about 30 percent of 2025 records, so these shares are provisional.

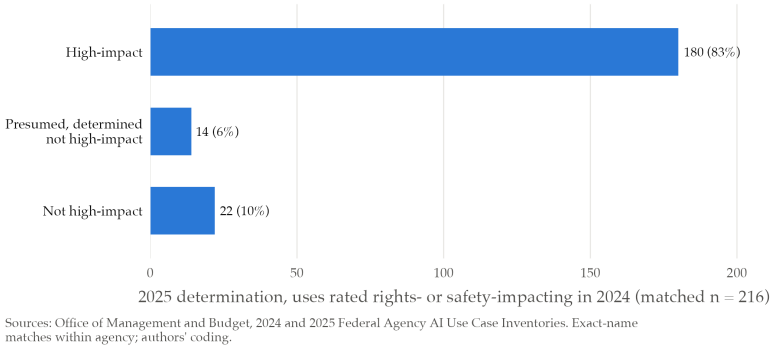


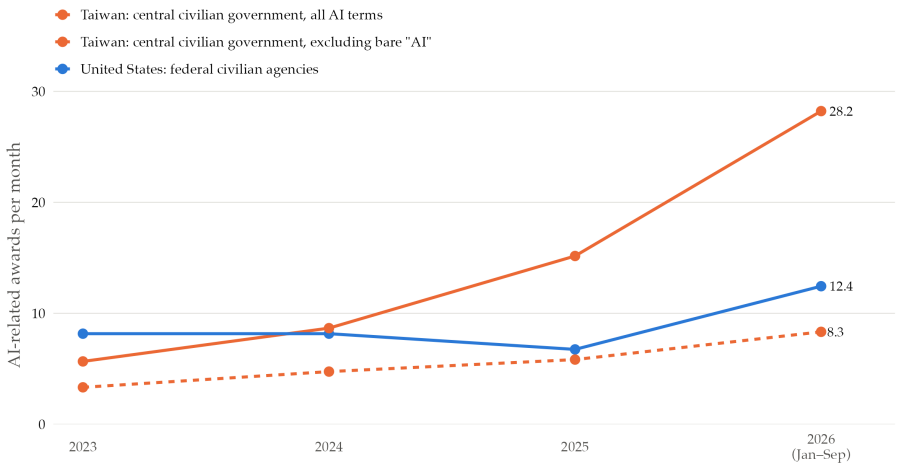
Figure 3: 2025 determination of uses rated rights- or safety-impacting in 2024, exact-name matches.

[Interpretation pending fuzzy matching. Report as a pattern in disclosure, not as evidence of bad faith: the 2025 definition differs from the 2024 categories, and some reclassification is expected.]

Preliminary Evidence: Procurement in Both Countries

Procurement records are the one source both governments publish in comparable form. We compare new AI-related awards from January 2023 to September 2026 for US federal civilian agencies (USAspending.gov, excluding the Department of Defense) and Taiwan's central civilian government (Government e-Procurement System; award notices only). Buyers in Taiwan are classified by level of govern-

ment from agency names. Because 2026 covers nine months, Figure 4 reports awards per month.



Sources: USAspending.gov contract awards (keyword search of award descriptions; Department of Defense excluded); Taiwan Government e-Procurement System via the g0v PCC API (title search; award notices of central civilian agencies). Counts are keyword matches, not yet classified for relevance. 2026 covers January–September.

Figure 4: AI-related procurement awards per month, US federal civilian agencies and Taiwan's central civilian government, January 2023 to September 2026. Counts are keyword matches, not yet classified for relevance.

Table 7: AI-related procurement awards by year and series (2026 = January–September).

Series	Year	Awards	Per month
United States: federal civilian agencies	2023	98	8.2
United States: federal civilian agencies	2024	98	8.2
United States: federal civilian agencies	2025	81	6.8
United States: federal civilian agencies	2026	112	12.4
Taiwan: central civilian government, all AI terms	2023	68	5.7
Taiwan: central civilian government, all AI terms	2024	104	8.7
Taiwan: central civilian government, all AI terms	2025	182	15.2
Taiwan: central civilian government, all AI terms	2026	254	28.2
Taiwan: central civilian government, excluding bare "AI"	2023	40	3.3
Taiwan: central civilian government, excluding bare "AI"	2024	57	4.8
Taiwan: central civilian government, excluding bare "AI"	2025	70	5.8

Table 7: AI-related procurement awards by year and series (2026 = January–September).

Series	Year	Awards	Per month
Taiwan: central civilian government, excluding bare "AI"	2026	75	8.3

AI-titled awards by Taiwan's central civilian agencies rose from about 6 a month in 2023 to 28 a month in 2026. Titles matched only by the bare term "AI" drive much of that rise, and they include hardware purchases. Excluding them, the increase is from about 3 to 8 a month. US federal civilian awards are steadier: about 8 a month in 2023 and 2024, 7 in 2025 and 13 in the first nine months of 2026.

Two cautions apply before these series are compared directly. The US search matches award descriptions and the Taiwanese search matches tender titles, so the two are not equally sensitive. Neither series has yet been classified for relevance; that step uses LLM-assisted coding validated on a hand-coded sample (Component A). Taiwanese purchasing is also widely distributed: central civilian agencies account for only about a third of AI-titled awards, with local governments and universities each contributing roughly a quarter. A comparison confined to the central level therefore understates how much AI purchasing in Taiwan happens beyond it, which bears directly on P4.

Taiwan: Agency Rules and Matched Cases

[To be written from Components C and D. Placeholder structure:]

- The 2023 generative AI guidelines as baseline: disclosure, verification, ban on drafting classified documents (Executive Yuan, 2023).
- Implementation of the AI Basic Act: risk assessments (July 2026), internal control rules (by mid-2027), civil-service AI competencies.
- Coded completeness of ministry rules by dimension; relation to prior digital capacity (P4).
- Public opinion context: in a May 2025 telephone survey, 64.8 percent of respondents said AI improves the quality of government services (Huang, 2025).

What Transfers, and What Does Not

[To be written after the matched cases. Expected argument: practices built on organizational routines (staff review, fallback, contractual audit access) travel; those built on jurisdiction-specific authority (executive memoranda, framework

legislation) do not. State limits: no country ranking; no causal claim about regimes.]

Conclusion

[To be written. The argument to make, if supported: adoption has outrun answerability in both systems, but in different places. The US documents its gap in public; Taiwan's gap is harder to observe. Democratic resilience depends less on the volume of AI adoption than on whether correction and challenge are built in before deployment.]

Appendix A: Coding Inventory Fields

Inventory field	Minimum practice (M-25-21)	Recoded status
`hi_testing_conducted`	Pre-deployment testing	Yes = established
`hi_assessment_completed`	AI impact assessment	Yes = established
`hi_independent_review`	Independent review	Internal, CAIO or board review = established
`hi_ongoing_monitoring`	Ongoing monitoring	Yes = established
`hi_training_established`	Operator training	Training established = established
`hi_failsafe_presence`	Fail-safe or human override	Yes = established; not applicable excluded
`hi_appeal_process`	Appeal or remedy	Established = established; not applicable or precluded by law excluded

"In progress" is coded separately. Blank fields are coded "not reported". Code: `analysis/01_inventory_descriptives.R`.

Appendix B: Codebook

[Codebook v1 to be drafted in November 2026: items, 0–2 scales, decision rules, Chinese and English examples, LLM prompt, reliability results.]

References

Alon-Barkat, Saar and Madalina Busuioc (2023), "Human–AI interactions in public sector decision making: "automation bias" and "selective adherence" to algorithmic advice." *Journal of Public Administration Research and Theory*, 33, 153–169.

Ansell, Christopher, Eva Sørensen, and Jacob Torfing (2023), "Public administration and politics meet turbulence: The search for robust governance responses." *Public Administration*, 101, 3–22.

Beach, Derek and Rasmus Brun Pedersen (2019), *Process-Tracing Methods: Foundations and Guidelines*, 2 edition. University of Michigan Press, Ann Arbor.

Boese, Vanessa A., Amanda B. Edgell, Sebastian Hellmeier, Seraphine F. Maerz, and Staffan I. Lindberg (2021), "How democracies prevail: Democratic resilience as a two-stage process." *Democratization*, 28, 885–907.

Boin, Arjen and Michel J. G. van Eeten (2013), "The resilient organization." *Public Management Review*, 15, 429–445.

Bovens, Mark (2007), "Analysing and assessing accountability: A conceptual framework." *European Law Journal*, 13, 447–468.

Bovens, Mark and Stavros Zouridis (2002), "From street-level to system-level bureaucracies: How information and communication technology is transforming administrative discretion and constitutional control." *Public Administration Review*, 62, 174–184.

Bullock, Justin B. (2019), "Artificial intelligence, discretion, and bureaucracy." *The American Review of Public Administration*, 49, 751–761.

Bürkner, Paul-Christian (2017), "brms: An R package for bayesian multilevel models using Stan." *Journal of Statistical Software*, 80, 1–28.

Bürkner, Paul-Christian and Matti Vuorre (2019), "Ordinal regression models in psychology: A tutorial." *Advances in Methods and Practices in Psychological Science*, 2, 77–101.

Busuioc, Madalina (2021), "Accountable artificial intelligence: Holding algorithms to account." *Public Administration Review*, 81, 825–836.

Chen, Yu-Che, Michael J. Ahn, and Yi-Fan Wang (2023), "Artificial intelligence and public values: Value impacts and governance in the public sector." *Sustainability*, 15, 4796.

Dolowitz, David P. and David Marsh (2000), "Learning from abroad: The role of policy transfer in contemporary policy-making." *Governance*, 13, 5–23.

Duit, Andreas (2016), "Resilience thinking: Lessons for public administration." *Public Administration*, 94, 364–380.

Engstrom, David Freeman, Daniel E. Ho, Catherine M. Sharkey, and Mariano-Florentino Cuéllar (2020), "Government by algorithm: Artificial intelligence in federal administrative agencies." Technical report, Administrative Conference of the United States, URL https://www.acus.gov/report/government-algorithm-artificial-intelligence-federal-administrative-agencies.

Executive Yuan (2023), "Guidelines for the use of generative AI by the executive yuan and its subordinate agencies (( ) AI)." Approved August 31, 2023.

Executive Yuan (2026), "Establishment of the national artificial intelligence strategy special committee." Cabinet announcement, May 21, 2026, URL https://www.ey.gov.tw/Page/9277F759E41CCD91/999df6da-6539-4941-913f-5125d684c58d.

Gelman, Andrew and Jennifer Hill (2007), *Data Analysis Using Regression and Multilevel/Hierarchical Models*. Cambridge University Press, Cambridge.

George, Alexander L. and Andrew Bennett (2005), *Case Studies and Theory Development in the Social Sciences*. MIT Press, Cambridge, MA.

Gilardi, Fabrizio, Meysam Alizadeh, and Maël Kubli (2023), "ChatGPT outperforms crowd workers for text-annotation tasks." *Proceedings of the National Academy of Sciences*, 120, e2305016120.

Grimmelikhuijsen, Stephan and Albert Meijer (2022), "Legitimacy of algorithmic decision-making: Six threats and the need for a calibrated institutional response." *Perspectives on Public Management and Governance*, 5, 232–242.

Hsiao, Yu-Tang, Shu-Yang Lin, Audrey Tang, Darshana Narayanan, and Claudina Sarahe (2018), "vTaiwan: An empirical study of open consultation process in Taiwan." Technical report, SocArXiv.

Huang, Tong-Yi (2025), "The public's perceptions and usage of AI in Taiwan." Presentation of telephone survey results (fieldwork May 19–27, 2025; n = 1,171), National Development Council, URL https://report.ndc.gov.tw/ReportFront/PageSystem/reportFileDownload/C11401217/002.

Krippendorff, Klaus (2018), *Content Analysis: An Introduction to Its Methodology*, 4 edition. SAGE, Thousand Oaks, CA.

Kuziemski, Maciej and Gianluca Misuraca (2020), "AI governance in the public sector: Three tales from the frontiers of automated decision-making in democratic settings." *Telecommunications Policy*, 44, 101976.

Lawrence, Christie, Isaac Cui, and Daniel E. Ho (2023), "The bureaucratic challenge to AI governance: An empirical assessment of implementation at U.S. federal agencies." In *Proceedings of the 2023 AAAI/ACM Conference on AI, Ethics, and Society*, ACM, Montréal, QC.

Legislative Yuan (2026), "Artificial intelligence basic act ( )." Passed third reading December 23, 2025; in force January 14, 2026.

Madan, Rohit and Mona Ashok (2023), "AI adoption and diffusion in public administration: A systematic literature review and future research agenda." *Government Information Quarterly*, 40, 101774.

Merkel, Wolfgang and Anna Lührmann (2021), "Resilience of democracies: Responses to illiberal and authoritarian challenges." *Democratization*, 28, 869–884.

Moore, Mark H. (1995), *Creating Public Value: Strategic Management in Government*. Harvard University Press, Cambridge, MA.

National Institute of Standards and Technology (2023), "Artificial intelligence risk management framework (AI RMF 1.0)." Technical Report NIST AI 100-1, U.S. Department of Commerce.

Neumann, Oliver, Katharina Guirguis, and Reto Steiner (2024), "Exploring artificial intelligence adoption in public organizations: A comparative case study." *Public Management Review*, 26, 114–141.

Office of Management and Budget (2024), "Advancing governance, innovation, and risk management for agency use of artificial intelligence." Memorandum M-24-10. Rescinded by M-25-21.

Office of Management and Budget (2025a), "2024 federal agency AI use case inventory." Dataset, GitHub repository ombegov/2024-Federal-AI-Use-Case-Inventory, URL https://github.com/ombegov/2024-Federal-AI-Use-Case-Inventory.

Office of Management and Budget (2025b), "Accelerating federal use of AI through innovation, governance, and public trust." Memorandum M-25-21, URL https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-21-Accelerating-Federal-Use-of-AI-through-Innovation-Governance-and-Public-Trust.pdf.

Office of Management and Budget (2025c), "Driving efficient acquisition of artificial intelligence in government." Memorandum M-25-22, URL https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Driving-Efficient-Acquisition-of-Artificial-Intelligence-in-Government.pdf.

Office of Management and Budget (2025d), "Increasing public trust in artificial intelligence through unbiased AI principles." Memorandum M-26-04, URL https://www.whitehouse.gov/wp-content/uploads/2025/12/M-26-04-Increasing-Public-Trust-in-Artificial-Intelligence-Through-Unbiased-AI-Principles-1.pdf.

Office of Management and Budget (2026), "2025 federal agency AI use case inventory." Dataset, GitHub repository ombegov/2025-Federal-Agency-AI-Use-Case-Inventory, URL https://github.com/ombegov/2025-Federal-Agency-AI-Use-Case-Inventory. Version of April 13, 2026.

Oxford Insights (2024), "Government AI readiness index 2024." Technical report, Oxford Insights, URL https://oxfordinsights.com/ai-readiness/ai-readiness-index/.

Raji, Inioluwa Deborah, Andrew Smart, Rebecca N. White, Margaret Mitchell, Timnit Gebru, Ben Hutchinson, Jamila Smith-Loud, Daniel Theron, and Parker Barnes (2020), "Closing the AI accountability gap: Defining an end-to-end framework for internal algorithmic auditing." In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 33–44, ACM.

Sabel, Charles F. and Jonathan Zeitlin (2012), "Experimentalist governance." In *The Oxford Handbook of Governance* (David Levi-Faur, ed.), 169–183, Oxford University Press, Oxford.

Schiff, Daniel S., Kaylyn Jackson Schiff, and Patrick Pierson (2022), "Assessing public value failure in government adoption of artificial intelligence." *Public Administration*, 100, 653–673.

Selten, Friso, Marcel Robeer, and Stephan Grimmelikhuijsen (2023), "'just like I thought': Street-level bureaucrats trust AI recommendations that align with their professional judgment." *Public Administration Review*, 83, 263–278.

U.S. Government Accountability Office (2023), "Artificial intelligence: Agencies have begun implementation but need to complete key requirements." Technical Report GAO-24-105980, U.S. Government Accountability Office, URL https://www.gao.gov/products/gao-24-105980.

Wieringa, Maranke (2020), "What to account for when accounting for algorithms: A systematic literature review on algorithmic accountability." In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 1–18, ACM.

Wirtz, Bernd W., Jan C. Weyerer, and Carolin Geyer (2019), "Artificial intelligence and the public sector—applications and challenges." *International Journal of Public Administration*, 42, 596–615.

Yeung, Karen (2018), "Algorithmic regulation: A critical interrogation." *Regulation & Governance*, 12, 505–523.

Young, Matthew M., Justin B. Bullock, and Jesse D. Lecy (2019), "Artificial discretion as a tool of governance: A framework for understanding the impact of artificial intelligence on public administration." *Perspectives on Public Management and Governance*, 2, 301–313.

Zuiderwijk, Anneke, Yu-Che Chen, and Fadi Salem (2021), "Implications of the use of artificial intelligence in public governance: A systematic literature review and a research agenda." *Government Information Quarterly*, 38, 101577.